

# Pandas

## What is Pandas?

- Pandas is a Python library used for working with data sets
- Pandas is a **powerful library** for **data analysis and manipulation** in Python. It provides **DataFrames** and **Series**, which make handling structured data easy.
- It has functions for analyzing, cleaning, exploring, and manipulating data.
- The name "Pandas" has a reference to both "Panel Data", and "Python Data Analysis" and was created by Wes McKinney in 2008.

## Installation of pandas:

To install pandas, follow these steps based on your setup:

### 1. Using pip (Recommended)

If you have Python installed, you can install pandas using pip:

**pip install pandas**

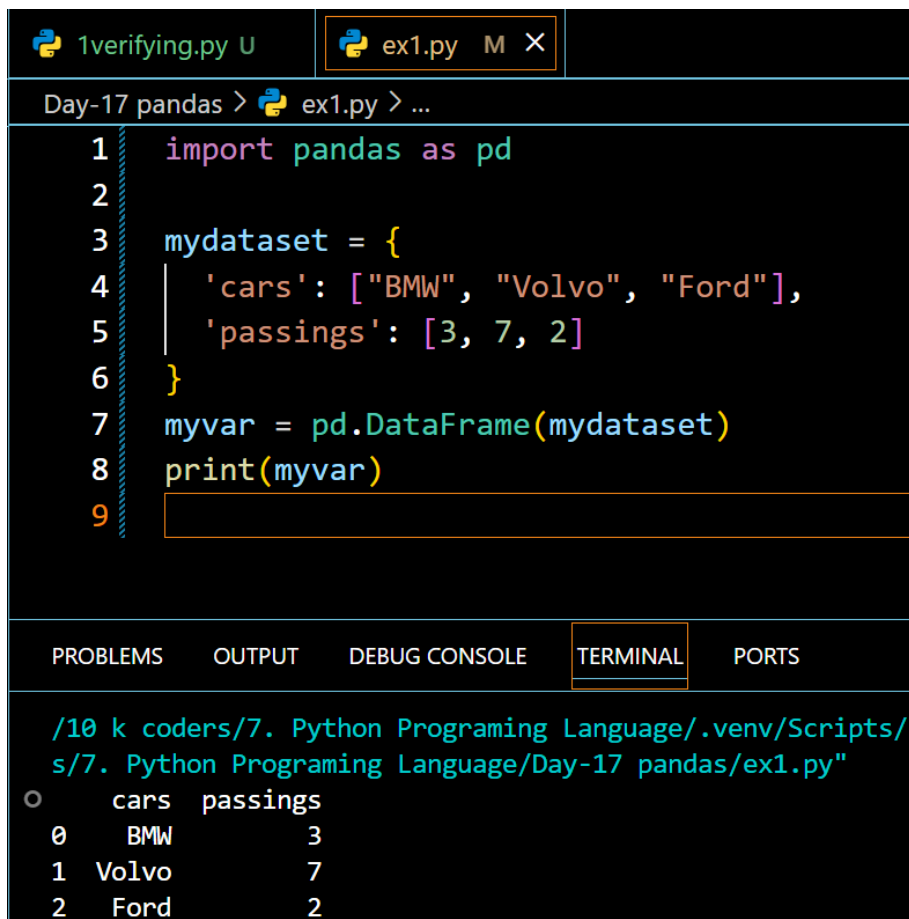
If you need to upgrade pandas, use:

**pip install --upgrade pandas**

### Verifying Installation

After installation, check if pandas is installed by running:

```
import pandas as pd
print(pd.__version__)
```



The screenshot shows a code editor with a file named `ex1.py` open. The code defines a dictionary `mydataset` with two keys: `'cars'` (values: "BMW", "Volvo", "Ford") and `'passings'` (values: 3, 7, 2). This dictionary is then used to create a `DataFrame` object `myvar`, which is printed to the console.

```
1 import pandas as pd
2
3 mydataset = {
4     'cars': ["BMW", "Volvo", "Ford"],
5     'passings': [3, 7, 2]
6 }
7 myvar = pd.DataFrame(mydataset)
8 print(myvar)
9
```

The terminal output shows the resulting DataFrame:

```
/10 k coders/7. Python Programing Language/.venv/Scripts/p
s/7. Python Programing Language/Day-17 pandas/ex1.py"
○   cars  passings
0   BMW         3
1  Volvo         7
2   Ford         2
```

# Pandas Series

## What is a Series?

A Pandas Series is like a column in a table.

It is a one-dimensional array holding data of any type.

```
Day-17 pandas > ex2_series.py > ...
1  import pandas as pd
2
3  a = [1, 7, 2]
4  myvar = pd.Series(a)
5  print(myvar)
6

PROBLEMS  OUTPUT  DEBUG CONSOLE  TERMINAL  PORTS

/10 k coders/7. Python Programing Language/.venv/Scripts/py
s/7. Python Programing Language/Day-17 pandas/ex2_series.py
0      1
1      7
2      2
dtype: int64
```

## Labels

If nothing else is specified, the values are labeled with their index number. First value has index 0, second value has index 1 etc.

This label can be used to access a specified value.

```
import pandas as pd

a = [1, 7, 2]

myvar = pd.Series(a)
print(myvar[1])

7
```

## Create Labels

With the index argument, you can name your own labels.

```
import pandas as pd

a = [1, 7, 2]

myvar = pd.Series(a, index = ["x", "y", "z"])

x      1
y      7
z      2
dtype: int64
```

## Key/Value Objects as Series

You can also use a key/value object, like a dictionary, when creating a Series.

```
import pandas as pd

calories = {"day1": 420, "day2": 380, "day3": 390}

myvar = pd.Series(calories)

print(myvar)
```

```
day1    420
day2    380
day3    390
dtype: int64
```

To select only some of the items in the dictionary, use the index argument and specify only the items you want to include in the Series.

```
import pandas as pd

calories = {"day1": 420, "day2": 380, "day3": 390}
myvar = pd.Series(calories, index = ["day1", "day2"])
print(myvar)
```

```
day1    420
day2    380
dtype: int64
```

## DataFrames

A Pandas DataFrame is a 2 dimensional data structure, like a 2 dimensional array, or a table with rows and columns.

```
import pandas as pd

data = {
    "calories": [420, 380, 390],
    "duration": [50, 40, 45]
}
df = pd.DataFrame(data)
print(df)
```

```
   calories  duration
0        420         50
1        380         40
2        390         45
```

## Locate Row

As you can see from the result above, the DataFrame is like a table with rows and columns. Pandas use the loc attribute to return one or more specified row(s)

```
import pandas as pd

data = {
    "calories": [420, 380, 390],
    "duration": [50, 40, 45]
}
#load data into a DataFrame object:
df = pd.DataFrame(data)
print(df.loc[0])
```

```
calories    420
duration     50
Name: 0, dtype: int64
```

## Example

Return row 0 and 1:

```
import pandas as pd

data = {
    "calories": [420, 380, 390],
    "duration": [50, 40, 45]
}
#load data into a DataFrame object:
df = pd.DataFrame(data)
print(df.loc[[0, 1]])
```

```
   calories  duration
0        420         50
1        380         40
```

## Named Indexes

With the index argument, you can name your own indexes.

### Example

Add a list of names to give each row a name:

```
import pandas as pd

data = {
    "calories": [420, 380, 390],
    "duration": [50, 40, 45]
}
df = pd.DataFrame(data, index = ["day1", "day2", "day3"])
print(df)
```

```
   calories  duration
day1        420         50
day2        380         40
day3        390         45
```

## Locate Named Indexes

Use the named index in the loc attribute to return the specified row(s).

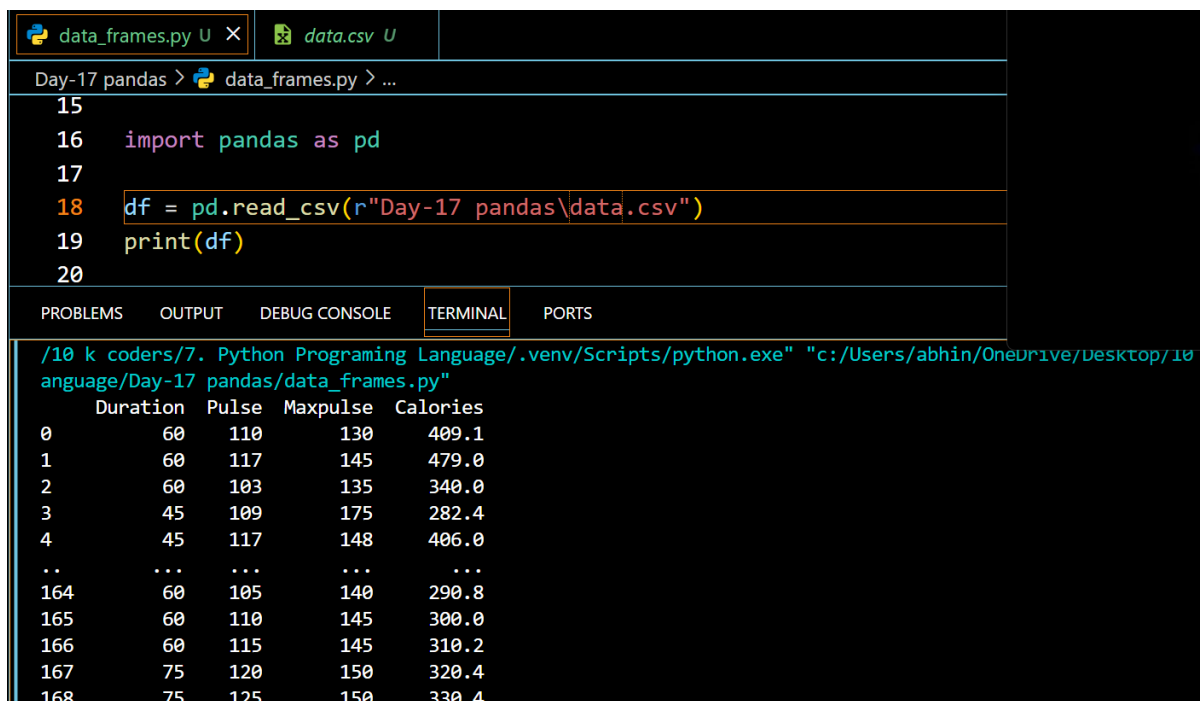
```
import pandas as pd

data = {
    "calories": [420, 380, 390],
    "duration": [50, 40, 45]
}
df = pd.DataFrame(data, index = ["day1", "day2", "day3"])
print(df.loc["day2"])
```

```
calories    380
duration     40
Name: day2, dtype: int64
```

## Read CSV Files

- A simple way to store big data sets is to use CSV files (comma separated files).
- CSV files contains plain text and is a well know format that can be read by everyone including Pandas.
- In our examples we will be using a CSV file called 'data.csv'.
- [Download data.csv](#). or [Open data.csv](#).



The screenshot shows a code editor with two tabs: 'data\_frames.py' and 'data.csv'. The 'data\_frames.py' tab is active, displaying the following code:

```
15
16 import pandas as pd
17
18 df = pd.read_csv(r"Day-17 pandas\data.csv")
19 print(df)
20
```

Below the code editor is a terminal window with the following output:

```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
/10 k coders/7. Python Programing Language/.venv/Scripts/python.exe "c:/Users/abhin/OneDrive/Desktop/10
anguange/Day-17 pandas/data_frames.py"
   Duration  Pulse  Maxpulse  Calories
0         60    110       130     409.1
1         60    117       145     479.0
2         60    103       135     340.0
3         45    109       175     282.4
4         45    117       148     406.0
..      ...    ...      ...      ...
164        60    105       140     290.8
165        60    110       145     300.0
166        60    115       145     310.2
167        75    120       150     320.4
168        75    125       150     330.4
```

```
Day-17 pandas > data_frames.py > ...
15
16 import pandas as pd
17
18 df = pd.read_csv(r"Day-17 pandas\data1.csv")
19 print(df.loc[1])
20
```

PROBLEMS OUTPUT DEBUG CONSOLE **TERMINAL** PORTS

```
PS C:\Users\abhin\OneDrive\Desktop\10 k coders\7. Python Programing Language>
/10 k coders/7. Python Programing Language/.venv/Scripts/python.exe" "c:/User
s/7. Python Programing Language/Day-17 pandas/data_frames.py"
Duration      60.0
Pulse         103.0
Maxpulse      135.0
Calories      340.0
Name: 1, dtype: float64
```

**Tip:** use `to_string()` to print the entire DataFrame.

```
import pandas as pd

df = pd.read_csv('data.csv')

print(df.to_string())
```

	Duration	Pulse	Maxpulse	Calories
0	60	110	130	409.1
1	60	117	145	479.0
2	60	103	135	340.0
3	45	109	175	282.4
4	45	117	148	406.0
5	60	102	127	300.5
6	60	110	136	374.0
7	45	104	134	253.3
8	30	109	133	195.1
9	60	98	124	269.0

## Read JSON

Big data sets are often stored, or extracted as JSON.

JSON is plain text, but has the format of an object, and is well known in the world of programming, including Pandas.

In our examples we will be using a JSON file called 'data.json'

```
import pandas as pd
df = pd.read_json('data.json')
print(df.to_string())
```

	Duration	Pulse	Maxpulse	Calories
0	60	110	130	409.1
1	60	117	145	479.0
2	60	103	135	340.0
3	45	109	175	282.4
4	45	117	148	406.0
5	60	102	127	300.5
6	60	110	136	374.0
7	45	104	134	253.3
8	30	109	133	195.1
9	60	98	124	269.0
10	60	103	147	329.3
11	60	100	120	250.7

## Basic Methods:

### Display the First Few Rows:

```
print(df.head()) # Default is 5 rows  
print(df.head(3)) # First 3 rows
```

### Display the Last Few Rows:

```
print(df.tail()) # Default is 5 rows
```

### Get DataFrame Information:

```
print(df.info())
```

### Get Summary Statistics:

```
print(df.describe())
```

### Get Column Names:

```
print(df.columns)
```

### Get Index Information:

```
print(df.index)
```

```
panda.py > ...  
38 data = [  
39     ["Alice", 25, "New York"],  
40     ["Bob", 30, "Los Angeles"],  
41     ["Charlie", 35, "Chicago"]  
42 ]  
43 df=pd.DataFrame(data)  
44  
45 # print(df.head())  
46 # print(df.tail())  
47 # print(df.info())  
48 print(df.describe())
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS powershell

None

(myenv1) PS C:\Users\laksh\OneDrive\Desktop\pandas> py panda.py

count	3.0
mean	30.0
std	5.0
min	25.0
25%	27.5
50%	30.0
75%	32.5

```
Welcome panda.py X  
panda.py > ...  
45 df=pd.DataFrame(data)  
46  
47 # print(df.head())  
48 # print(df.tail())  
49 # print(df.info())  
50 # print(df.describe())  
51 print(df.columns)  
52 print(df.index)
```

```
(myenv1) PS C:\Users\laksh\OneDrive\Desktop\pandas> py panda.py  
RangeIndex(start=0, stop=3, step=1)  
RangeIndex(start=0, stop=5, step=1)  
(myenv1) PS C:\Users\laksh\OneDrive\Desktop\pandas>
```

## Selecting methods

### Selecting a Single Column:

```
print(df["Name"])
```

### Selecting Multiple Columns:

```
print(df[["Name", "Age"]])
```

### Selecting Rows by Index:

```
print(df.loc[0]) # Select the first row
```

### Selecting Rows and Columns:

```
print(df.loc[0, "Name"]) # Row 0, Column "Name"
```

### Selecting Using `iloc` (Position-Based Indexing):

```
print(df.iloc[0, 1]) # First row, second column
```

## 4. Filtering Data

### Filtering Rows Based on a Condition:

```
filtered_df = df[df["Age"] > 30]
print(filtered_df)
```

### Filtering Rows Using Multiple Conditions:

```
filtered_df = df[(df["Age"] > 25) & (df["City"] == "Chicago")]
print(filtered_df)
```

## 5. Modifying Data

### Adding a New Column:

```
df["Salary"] = [50000, 60000, 70000]
print(df)
```

### Updating a Column:

```
df["Age"] = df["Age"] + 5
print(df)
```

### Deleting a Column:

```
df = df.drop("Salary", axis=1)
print(df)
```

```
panda.py > ...
43 df=pd.DataFrame(data)
44
45 # print(df["Name"])
46 # print(df["City"])
47
48 df["Salary"]=[10000,20000,30000,0]
49 print(df[df["Salary"]==10000])
50

PS C:\Users\laksh\OneDrive\Desktop\pandas> py panda.py
   Name  Age  City  Salary
0  Alice   45  New York  10000

PS C:\Users\laksh\OneDrive\Desktop\pandas>
```



```
panda.py x
panda.py > ...
47
48 df["Salary"]=[10000,20000,30000,0]
49 # print(df[df["Salary"]==10000]=df[df["Salary"]+5000])
50
51
52 # df["Salary"]=df["Salary"]+5000
53 # print(df)
54 df=df.drop("Salary",axis=1)
55 ## axis is used to define whether to delete a row or column (axis=0 is for rows) (axis=1) for columns
56 print(df)
57
58
59
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

2 Charlie 35 Chicago 30000  
3 Jack 50 New York 0  
PS C:\Users\laksh\OneDrive\Desktop\pandas> py panda.py

	Name	Age	City
0	Alice	45	New York
1	Bob	30	Los Angeles

## 6. Sorting Data

### Sorting by a Column:

```
sorted_df = df.sort_values("Age")  
print(sorted_df)
```

### Sorting in Descending Order:

```
sorted_df = df.sort_values("Age", ascending=False)  
print(sorted_df)
```

### Sorting by Multiple Columns:

```
sorted_df = df.sort_values(["Age", "City"])  
print(sorted_df)
```

## 7. Handling Missing Data

### Detecting Missing Values:

```
print(df.isnull())
```

### Filling Missing Values:

```
df["Age"] = df["Age"].fillna(30)  
print(df)
```

### Dropping Rows with Missing Values:

```
df = df.dropna()  
print(df)
```

```
62 df=pd.DataFrame([{"a":10,"b":20}, {"a":10,"b":20,"c":40}])
63 print(df)
64
65
```

PROBLEMS OUTPUT DEBUG CONSOLE **TERMINAL** PORTS

```
1 Bob 30 Los Angeles 20000
PS C:\Users\laksh\OneDrive\Desktop\pandas> py panda.py
a b c
0 10 20 NaN
1 10 20 40.0
PS C:\Users\laksh\OneDrive\Desktop\pandas>
```

```
61
62 df=pd.DataFrame([{"a":10,"b":20}, {"a":10,"b":20,"c":40}])
63 # print(df.isnull())
64 print(df["c"].fillna(0))
65
66
```

PROBLEMS OUTPUT DEBUG CONSOLE **TERMINAL** PORTS

```
Name: c, dtype: float64
PS C:\Users\laksh\OneDrive\Desktop\pandas> py panda.py
0 0.0
1 40.0
Name: c, dtype: float64
PS C:\Users\laksh\OneDrive\Desktop\pandas>
```

```
62 df=pd.DataFrame([{"a":10,"b":20}, {"a":10,"b":20,"c":40}])
63 print(df.isnull())
64 print(df.dropna())
65
66
```

PROBLEMS OUTPUT DEBUG CONSOLE **TERMINAL** PORTS

```
a b c
0 False False True
1 False False False
a b c
1 10 20 40.0
```

## 8. Grouping and Aggregating

### Grouping Data by a Column:

```
grouped = df.groupby("City").mean()
print(grouped)
```

### Aggregating Data:

```
aggregated = df.groupby("City").agg({"Age": "max", "Salary": "sum"})
print(aggregated)
```

## 9. Combining DataFrames

### Concatenating DataFrames:

```
df1 = pd.DataFrame({"A": [1, 2], "B": [3, 4]})  
df2 = pd.DataFrame({"A": [5, 6], "B": [7, 8]})  
combined = pd.concat([df1, df2])  
print(combined)
```

### Merging DataFrames:

```
df1 = pd.DataFrame({"Name": ["Alice", "Bob"], "Age": [25, 30]})  
df2 = pd.DataFrame({"Name": ["Alice", "Bob"], "City": ["New York", "Chicago"]})  
merged = pd.merge(df1, df2, on="Name")  
print(merged)
```

```
68 df1 = pd.DataFrame({"A": [1, 2], "B": [3, 4]})  
69 df2 = pd.DataFrame({"C": [5, 6], "B": [7, 8]})  
70 combined = pd.concat([df2, df1])  
71 print(combined)
```

	C	B	A
0	5.0	7	NaN
1	6.0	8	NaN
0	NaN	3	1.0
1	NaN	4	2.0

PS C:\Users\laksh\OneDrive\Desktop\pandas>

### Writing to a CSV file:

df.to\_csv('output.csv', index=False)

```
41 "City": ["York", "Los Angeles", "Chicago", "New York"]  
42 }  
43 }  
44  
45 # print(df["Name"])  
46 # print(df["City"])  
47  
48 df["Salary"]=[10000,20000,30000,10000]  
49 # print(df[df["Salary"]==10000]=df[df["Salary"]+5000])  
50 df.to_csv("output.csv", index=False)  
51
```

### Removing Duplicates

To discover duplicates, we can use the duplicated() method.

The duplicated() method returns a Boolean values for each row:

#### Example

Returns **True** for every row that is a duplicate, otherwise **False** :

```
print(df.duplicated())
```

#### Removing Duplicates

To remove duplicates, use the **drop\_duplicates()** method.

#### Example

Remove all duplicates:

```
df.drop_duplicates(inplace = True)
```

# Pandas - Data Correlations

## Finding Relationships

A great aspect of the Pandas module is the `corr()` method.

The `corr()` method calculates the relationship between each column in your data set.

The examples in this page uses a CSV file called: 'data.csv'.

[Download data.csv](#) or [Open data.csv](#)

```
import pandas as pd
df = pd.read_csv('data.csv')
print(df.corr())
```

	Duration	Pulse	Maxpulse	Calories
Duration	1.000000	-0.059452	-0.250033	0.344341
Pulse	-0.059452	1.000000	0.269672	0.481791
Maxpulse	-0.250033	0.269672	1.000000	0.335392
Calories	0.344341	0.481791	0.335392	1.000000

## Result Explained

- The Result of the `corr()` method is a table with a lot of numbers that represents how well the relationship is between two columns.
- The number varies from -1 to 1.
- 1 means that there is a 1 to 1 relationship (a perfect correlation), and for this data set, each time a value went up in the first column, the other one went up as well.
- 0.9 is also a good relationship, and if you increase one value, the other will probably increase as well.
- -0.9 would be just as good relationship as 0.9, but if you increase one value, the other will probably go down.
- 0.2 means NOT a good relationship, meaning that if one value goes up does not mean that the other will.

## Plotting

- Pandas uses the `plot()` method to create diagrams.
- We can use Pyplot, a submodule of the Matplotlib library to visualize the diagram on the screen.
- Read more about [ploting](#)